



Assignment 3

Hand in this assignment until Friday, November 14th 2025, 12:00 pm at the latest.

Running out of ideas?

Are you hitting a roadblock? Are some of the exercises unclear? Do you just need that one hint to get the ball rolling? Refer to the #forum channel on our Discord server—maybe you'll find just the help you need.

Exam-style Exercises

Exercises marked with are similar in style to those you will find in the exam. You can use these to hone your expectations and gauge your skills.

Task 1: So similar, yet so different

Create an instance for table `r` with schema `r(a int, b int)` such that the two queries Q1 and Q2 compute different results.

Query Q1: `WHERE`-clause

```
1 SELECT r.a, count(*) AS c
2 FROM r AS r
3 WHERE r.b <> 3
4 GROUP BY r.a;
```

Query Q2: `HAVING`-clause

```
1 SELECT r.a, count(*) AS c
2 FROM r AS r
3 GROUP BY r.a
4 HAVING bool_and(r.b <> 3);
```

Note

Query Q2 uses the aggregate function `bool_and`. Read about it in the [DuckDB documentation](#) .

Task 2: Production Steps

Consider table `production` which tracks the progress of items currently in production. The production process of these items occurs in multiple steps. If a step has been *completed* for an item, we store the completion time and date. If the step is still *in progress*, we write `NULL` into column `completion` holds `NULL`.

```
1 CREATE TABLE production (
2     item      char(20) NOT NULL,
3     step      int      NOT NULL,
4     completion timestamp, -- NULL ≡ incomplete
5     PRIMARY KEY (item, step)
6 );
```

To the right, you can see an example instance of `production`. Notice that step 2 of AT-AT production is not complete yet. Similarly for DS II production, step 1 is still pending, but steps 2 and 3 are already completed.

- A Formulate a SQL query that lists the names of all items for which all production steps are complete. Avoid duplicate item names. For the example instance to the right, we expect the following result:

item	step	completion
TIE	1	1977-03-02 04:12:00
TIE	2	1977-12-29 05:55:00
AT-AT	1	1978-01-03 14:12:00
AT-AT	2	NULL
DSII	1	NULL
DSII	2	1979-05-26 20:05:00
DSII	3	1979-04-04 17:12:00

complete
TIE

Task 3: Measurements

You are handed a note by a physicist, who kindly asks you to take a look at the measurements $m(t)$ (taken at time t) they intend to store inside a RDBMS of your choice (DuckDB, of course). They explain that at each timestamp t , one or more measurements $m(t)$ have been recorded.

Measurements Note	
t	$m(t)$
1.0	1.0, 3.0, 5.0, 5.0
2.5	0.8, 2.0
4.0	0.5
5.5	3.0
8.0	2.0, 6.0, 8.0
10.5	1.0, 3.0, 8.0

Write SQL queries based on the measurements table (provided in `measurements.sql`) for the following tasks. Assume that, in the future, more measurements may be recorded.

- [A] Compute a *one-column* table with the global maximum of the measurements $m(t)$.
- [B] Compute a *two-column* table that lists each time t and the average of its measurements $m(t)$.
- [C] Compute a *two-column* table that lists the average of all measurements $m(t)$ in each timeframe $[0.0 - 5.0)$, $[5.0 - 10.0)$, $[10.0 - 15.0)$, Be aware that there might be more measurements in different timeframes in the future.
- [D] Find all times t (there may be more than one) at which the global maximum was recorded. The result is a *two-column* table that lists time t and the global maximum.