

# Advanced SQL

---

②

The Core of SQL

Winter 2025/26

Torsten Grust  
Universität Tübingen, Germany

# 1 | The Core of SQL

---

- Let us recollect the **core constructs of SQL**, synchronize notation, and introduce query conventions.
- If you need to refresh your SQL memory, consider
  - the notes for [TaDa](#) (Chapters 6 to 8)
  - the [DuckDB web](#) (Documentation > SQL)
- We will significantly expand on this base SQL vocabulary during the semester.

## Sample Table

---

Table **T** serves as a common “playground” for the upcoming SQL queries:

<u>a</u>	<u>b</u>	<u>c</u>	<u>d</u>
1	'x'	true	10
2	'y'	true	40
3	'x'	false	30
4	'y'	false	20
5	'x'	true	NULL

Table **T**

```
CREATE TABLE T (a int PRIMARY KEY, -- implies NOT NULL
                  b text,           -- here: char(1)
                  c boolean,
                  d int);
```

## 2 : Row Variables

---

- Iterate over all rows of table **T** (in *some* order: bag semantics), bind **row variable** **t** to current row:

```
SELECT t          -- 2 t is bound to current row of T
FROM   T AS t     -- 1 bind/introduce t
```

- If you omit **AS t** in the **FROM** clause, a row variable **T** (generally: **AS <table name>**) will be implicitly introduced.
- This course: always explicitly introduce/name row variables for disambiguation, clarity, readability.

## Row Values

---

```
SELECT t          -- 2 t is bound to current row of T
FROM   T AS t     -- 1 bind/introduce t
```

- Row variable `t` is iteratively bound to **row values** whose field values and types are determined by the rows of table `T`:

<b>field names:</b>	a	b	c	d	
	↓	↓	↓	↓	
	t ≡ {a:5,	b:'x',	c:true,	d:NULL}	} <b>row values</b> (structs)
	t ≡ {a:1,	b:'x',	c:true,	d:10}	
	⋮				
	t ≡ {a:2,	b:'y',	c:true,	d:40}	
	↑	↑	↑	↑	
<b>field types:</b>	<u>int</u>	<u>text</u>	<u>boolean</u>	<u>int</u>	

## Row Types (Struct Types)

---

- All rows  $t$  in table  $T$  have **row type** `row(a int, b text, c boolean, d int)`<sup>1</sup>. Abbreviated as `t :: row(...)`. Read symbol `::` as “*has type*.”
- A row type  $\tau$  can also be explicitly defined via

```
CREATE TYPE  $\tau$  AS row(a int, b text, c boolean, d int);
```

- New row type  $\tau$  is then usable like any builtin type:


```
TABLE duckdb_types;      -- inspect DuckDB's catalog
```

<sup>1</sup> In DuckDB, `struct` is a synonym for `row`: `struct(a int, b text, c boolean, d int)`.

## Row Field Access and \* (“Star”)

---

- Named **field access** uses dot notation. Assume  $t :: \tau$  and binding  $t \equiv \{a:5, b:'x', c:true, d:NULL\}$ , then:
  - $t.b$  evaluates to `'x'` (of type `text`),
  - $t.d$  evaluates to `NULL` (of type `int`).
- Field names are *not* first-class in SQL and must be provided verbatim (i.e., may *not* be computed).
- Notation  $t.*$  abbreviates  $t.a, t.b, t.c, t.d$  in contexts where this makes sense.<sup>2</sup>

<sup>2</sup>  $t.*$  is most often used in `SELECT` clauses. DuckDB adds lots of syntactic sugar when it comes to  $*$  in the `SELECT` clause, see [Star Expression](#)  in the DuckDB documentation.

## Row Comparisons

---

- **Row comparisons** between rows  $t_1$ ,  $t_2$  are performed field-by-field and lexicographically (provided that the field types match). Assume  $t_1 :: \tau$ ,  $t_2 :: \tau$ :
  - $t_1 = t_2 \iff t_1.a = t_2.a \text{ AND } \dots \text{ AND } t_1.d = t_2.d$
  - $t_1 < t_2 \iff$   
 $t_1.a < t_2.a \text{ OR } (t_1.a = t_2.a \text{ AND } t_1.b < t_2.b) \text{ OR } \dots$
- Thus (here:  $\tau \equiv \text{row}(a \text{ int}, b \text{ text})$ ):
  - $\{a:1, b:'z'\} < \{a:2, b:'x'\}$  and
  - $\{a:2, b:'z'\} > \{a:2, b:'x'\}$ .



### 3 | The **SELECT** Clause

---

A **SELECT** clause evaluates  $n$  expressions  $e_1, \dots, e_n$ :

```
SELECT  $e_1$  AS  $c_1$ , ...,  $e_n$  AS  $c_n$ 
```

- Creates  $n$  columns named  $c_1, \dots, c_n$  (case-insensitive).
- In absence of **AS**  $c_i$ , DuckDB derives a column name from  $e_i$ . Almost always this is questionable practice ☹️.
- Thus: explicitly use **AS** to name columns unless a name can be sensibly derived from  $e_i$  (e.g., when  $e_i \equiv t.a$ ).
- If a column/table names contains whitespace/symbols/keywords: wrap in quotes " $c_i$ ".

## Standalone **SELECT**

---

- If query  $Q$  generates  $n$  row bindings, **SELECT** is evaluated  $n$  times to emit  $n$  rows (but see *aggregates* below).
- A standalone **SELECT** (no **FROM** clause) is evaluated exactly once and emits a single row:

```
SELECT 1 + 41 AS "The Answer", 'Gla' || 'DOS' AS Portal;
```

The Answer	Portal
42	GlaDOS

## 4 : Literal Tables (**VALUES**)

---

A **VALUES** clause constructs a **transient table** from a **list of  $n$  provided rows**:

```
VALUES ( $e_{11}, \dots, e_{1k}$ ), ..., ( $e_{n1}, \dots, e_{nk}$ )
```

- If  $n > 1$ , the rows must agree in arity  $k$  and field types (first row value is used to infer and determine types).
- **VALUES** automatically assigns column names "**col< $i$ >**". Use column aliasing to assign names (see **FROM** below).
- **Orthogonality:** a **VALUES** clause (in parentheses **(...)**) may be used anywhere a SQL query expects a table.

## 5 : Generating Row Variable Bindings (**FROM**)

---

A **FROM** clause expects a set of tables  $T_i$  and successively binds the row variables  $t_i$  to the tables' rows:

```
SELECT ... -- 2
FROM   T1 AS t1, ..., Tn AS tn -- 1
```

- The  $T_i$  may be table names or SQL queries computing tables (enclosed in (...)).
- If you need to rename the columns of  $T_i$  (recall the **VALUES** clause), use **column aliasing** on all (or only the first  $k$  ☹) columns:

$T_i \text{ AS } t_i(\underline{c_{i.1}}, \dots, \underline{c_{i.k}})$

## FROM Computes Cartesian Products

---

```
SELECT ...  
FROM     $T_1$  AS  $t_1$ , ...,  $T_n$  AS  $t_n$ 
```

- This **FROM** clause generates  $|T_1| \times \dots \times |T_n|$  bindings.  
Semantics: compute the **Cartesian product**  $T_1 \times \dots \times T_n$ , draw the bindings for the  $t_i$  from this product.
- **FROM** operates over a set of tables (',' is associative and commutative).
- In particular, row variable  $t_i$  is *not* in scope in the table subqueries  $T_{i+1}, \dots, T_n$  (but see **LATERAL**).

## 6 : WHERE Discards Row Bindings

---

A **WHERE** clause introduces a predicate  $p$  that is evaluated under all row variable bindings generated by **FROM**:

<b>SELECT</b>	...	--	3
<b>FROM</b>	$T_1$ <b>AS</b> $t_1$ , ..., $T_n$ <b>AS</b> $t_n$	--	1
<b>WHERE</b>	$p$	--	2

- All row variables  $t_i$  are in scope in predicate  $p$ .
- Only bindings that yield  $p = \text{true}$  are passed on.<sup>3</sup>
- Absence of a **WHERE** clause is interpreted as **WHERE true**.

<sup>3</sup> If  $p$  evaluates to **NULL**  $\neq$  **true**, the binding is discarded.

## 7 : Compositionality: Subqueries Instead of Values

---

*“The meaning of a complex expression is determined by the meanings of constituent expressions.”*

—Principle of Compositionality

With the advent of the SQL-92 and SQL:1999 standards, SQL has gained in **compositionality** and **orthogonality**:

- Wherever a (tabular or scalar) value *v* is required, a SQL expression in *(...)*—a **subquery**—may be used to compute *v*.
- Subqueries nest to arbitrary depth.


## Scalar Subqueries: Atomic Values

---

A SQL query that computes a **single-row, single-column table** (column name  $\square$  irrelevant) may be used in place of an atomic value  $v$ :



In a **scalar subquery**...

- ... an empty table is interpreted as **NULL**,
- ... a table with  $> 1$  columns yields a **compile-time error**,
- ... a table with  $> 1$  rows yields an **error at runtime**. <sup>4</sup>

<sup>4</sup> With DuckDB's configuration option `scalar_subquery_error_on_multiple_rows` set to `false`, a *random* value  $v$  will be selected from the table. This merely hides a potential source of errors. Avoid.



## Scalar Subqueries: Atomic Values

---

generate single column

↓

```
SELECT 2 + (SELECT t.d AS _
             FROM   T AS t
             WHERE  t.a = 2) AS "The Answer"
```

equality predicate on key column,  
will yield  $\leq 1$  rows


- **Compile-time error** for `SELECT t.a, t.d` ( $> 1$  column)
- **Runtime error** for `WHERE t.a > 2` ( $> 1$  row)
- Subquery yields `NULL`: `WHERE t.a = 0`
- `AS _` assigns a “*don't care*” column name—this is a case where column naming is obsolete and adds nothing.

## Row Variable Scoping


---

Subqueries may **refer to any row variable  $t$  bound in their enclosing queries** (up to the top-level query):

enclosing query



enclosing query



- **NB.** Seen from inside the subquery—*i.e.*, inside the `(...)`—row variable  $t$  is *free* (binding provided externally).

## Subqueries, Free Row Variables, Correlation

---

- If  $t$  is free in subquery  $q$ , we may understand the subquery as a function  $q(t)$ : you supply a value for  $t$ , I will compute the (tabular) value of  $q$ :

<pre> SELECT t1.* FROM   T AS t1 WHERE  t1.b &lt;&gt; (SELECT t2.b                 FROM   T AS t2                 WHERE  t1.a = t2.a) </pre>	}	<p>evaluated 5 times under t1 bindings:</p> <p>t1 ≡ (1, ...)</p> <p>t1 ≡ (2, ...)</p> <p>t1 ≡ (3, ...)</p> <p>t1 ≡ (4, ...)</p> <p>t1 ≡ (5, ...)</p>
<p>↑</p> <p>free</p>		

- Subqueries featuring free variables are also known as **correlated**.

## 8 : Row Ordering (**ORDER BY**)

---

SQL tables are **unordered bags** of rows, but rows may be **locally ordered** for result display or positional access:

<b>SELECT</b> ...	--	3
<b>FROM</b> ...	--	1
<b>WHERE</b> ...	--	2
<b>ORDER BY</b> $e_1, \dots, e_n$	--	4

- The order of the  $e_i$  matters: sort order is determined lexicographically with  $e_1$  being the major criterion.
- Sort criteria  $e_i$  are expressions that may refer to column names in the **SELECT** clause (evaluated after **SELECT**).

**SELECT t.\* FROM T AS t ...**

---

a	b	c	d
5	'x'	true	NULL
1	'x'	true	10
4	'y'	false	20
3	'x'	false	30
2	'y'	true	40

... ORDER BY t.d ASC NULLS FIRST

a	b	c	d
4	'y'	false	20
2	'y'	true	40
3	'x'	false	30
1	'x'	true	10
5	'x'	true	NULL

... ORDER BY t.b DESC, t.c

- Note: **ASC** (ascending) is default. **NULL** is larger than any non-**NULL** value. Ties: order is implementation-dependent.

## Row Order is Local Only

---

**ORDER BY** establishes a well-defined row order that is **local** to the current (sub)query:

may yield rows in any order

↓

```
SELECT t1.*  
FROM (SELECT t2.*  
      FROM T AS t2  
      ORDER BY t2.a) AS t1;
```

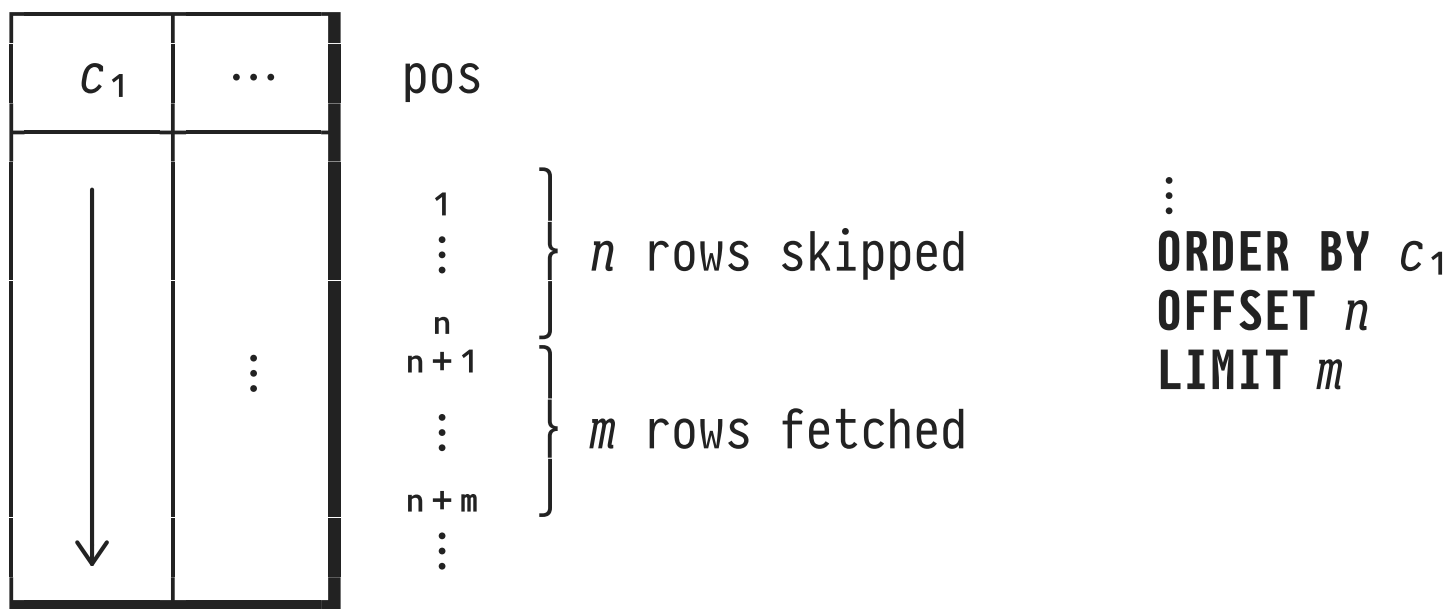
} guaranteed **row order**  
inside the subquery **only**

- ⚠ Never rely on row orders that appear consistent across runs—may vary between DBMSs, presence of indexes, etc.
- **Q:** What, then, is such local row order good for?

## Positional Access to Rows

---

Once row order has been established it makes sense to “skip to the  $n^{\text{th}}$  row” or “fetch the *next*  $m$  rows.”




- **OFFSET 0**: read from the start. **LIMIT ALL**: fetch all rows.
- Alternative **LIMIT** syntax: **FETCH [FIRST|NEXT]  $m$  ROWS ONLY**.

## 9 : Identify Particular Rows Among Peers (**DISTINCT ON**)

---

Extract the **first row among a group of equivalent rows**:

	prefix of <b>ORDER BY</b> clause	
	$\underbrace{\hspace{10em}}$	
<b>SELECT DISTINCT ON</b>	<b>4</b> $(e_1, \dots, e_n)$ $c_1, \dots, c_k$	-- 2
<b>FROM</b>	$\dots$	-- 1
<b>ORDER BY</b>	$e_1, \dots, e_n, e_{n+1}, \dots, e_m$	-- 3

1. Sort rows in  $e_1, \dots, e_n, e_{n+1}, \dots, e_m$  order.
  2. Rows with identical  $e_1, \dots, e_n$  values form one **group**.
  3. From each of these groups, pick **the first row** in  $e_{n+1}, \dots, e_m$  order.
-  Without **ORDER BY**, step 3 picks *any* row in each group.



## **DISTINCT ON:** Group, Then Pick First in Each Group

---

**SELECT DISTINCT ON ( $A_1$ ) ...**      -- For each  $A_1$ , pick the row ...  
**FROM** ...  
**ORDER BY  $A_1$ ,  $A_2$  DESC**      -- ... with the largest  $A_2$

	$A_1$	$A_2$	...
	$\vdots$	$\vdots$	$\vdots$
group {	$x_i$	$y_{i1}$	...
	$x_i$	$\vdots$	$\vdots$
group {	$x_j$	$y_{j1}$	...
	$x_j$	$\vdots$	$\vdots$
	$x_j$	$\vdots$	$\vdots$
	$\vdots$	$\vdots$	$\vdots$

← pick  
 } discard

← pick  
 } discard

## **DISTINCT:** Table-Wide Duplicate Removal

---

Keep only a single row from each group of **duplicates**:

```
SELECT DISTINCT ③  c1, ..., ck      -- ②
FROM             ...                -- ①
```

- True duplicate removal: rows are considered identical if they agree on **all** *k* columns *c<sub>i</sub>*.<sup>5</sup>
- Row order is irrelevant. **DISTINCT** returns a *set of rows*.
- May use **SELECT ALL ...** to explicitly document that a query is expected to return duplicate rows.

<sup>5</sup> This is equivalent to **SELECT DISTINCT ON (c<sub>1</sub>,...,c<sub>k</sub>) c<sub>1</sub>,...,c<sub>k</sub> FROM ...**.

## 10 : Summarizing Values: Aggregates

---

**Aggregate functions** (short: **aggregates**) reduce a *collection* of values to a *single* value (think summation, maximum).

- Simplest form: *collection*  $\equiv$  entire table:

```
SELECT agg1(e1) AS c1, ..., aggn(en) AS cn  
FROM ...
```

- Reduction of input rows: result table will have **one row**.
- Cannot mix aggregates with non-aggregate expression *e* in a **SELECT** clause:<sup>6</sup> which value of *e* should we pick?

<sup>6</sup> But see **GROUP BY** later on.

## Aggregate Functions: Semantics

---

```
SELECT agg(e) AS c  -- e will typically refer to t
FROM   T AS t      -- range over entire table T
```

- Aggregate *agg* defined by triple  $(\phi^{agg}, z^{agg}, \oplus^{agg})$ :
  - $\phi^{agg}$  (*empty*): aggregate of the empty value collection
  - $z^{agg}$  (*zero*): aggregate value initialiser
  - $\oplus^{agg}$  (*merge*): add value to existing aggregate

```
a ←  $\phi^{agg}$            -- a will be aggregate value
for t in T             -- iterate over all rows of T
| x ← e(t)             -- value to be aggregated
| if x ≠ NULL          -- aggregates ignore NULL values (*)
| | if a =  $\phi^{agg}$       -- once we see first non-NULL value:
| | | a ←  $z^{agg}$         -- initialize aggregate
| | | a ←  $\oplus^{agg}(a, x)$  -- maintain running aggregate
```

## Aggregate Functions: Semantics

---

Aggregate <i>agg</i>	$\phi^{agg}$	$z^{agg}$	$\oplus^{agg}(a, x)$
count	0	0	$a + 1$
sum	NULL <sup>7</sup>	0	$a + x$
avg <sup>8</sup>	NULL	$\langle 0, 0 \rangle$	$\langle a.1 + x, a.2 + 1 \rangle$
max	NULL	$-\infty$	$\max_2(a, x)$
min	NULL	$+\infty$	$\min_2(a, x)$
bool_and	NULL	true	$a \wedge x$
bool_or	NULL	false	$a \vee x$
⋮	⋮	⋮	⋮

- The special form **count(\*)** counts rows regardless of their fields' contents (NULL, in particular).

<sup>7</sup> If you think “*this is wrong*,” we’re two already. Possible upside: **sum** differentiates between summation over an empty collection vs. a collection of all 0s.

<sup>8</sup> Returns  $a.1 / a.2$  as final aggregate value.

## Aggregate Functions on Table T

---

```

SELECT count(*)           AS "#rows",
       count(t.d)         AS "#d",
       sum(t.d)           AS "Σd",
       max(t.b)           AS "max(b)",
       bool_and(t.c)       AS "∀c",
       bool_or(t.d = 30)  AS "∃d=30"
FROM   T AS t
WHERE  p

```

#rows	#d	Σd	max(b)	∀c	∃d=30
5	4	100	'y'	false	true

$p \equiv \text{true}$

#rows	#d	Σd	max(b)	∀c	∃d=30
0	0	NULL	NULL	NULL	NULL

$p \equiv \text{false}$

## Ordered Aggregates

---

- For most aggregates *agg*, merge  $\oplus^{agg}$  is commutative (and associative): row order does not matter.
- **Order-sensitive aggregates** admit a trailing `ORDER BY  $e_1, \dots, e_n$`  argument that defines row order:<sup>9</sup>

```
--          cast to text      separator string
--
SELECT string_agg(t.a :: text, ',' ORDER BY t.d) AS "all a"
FROM   T AS t
```

**all a**

'1,4,3,2,5'

<sup>9</sup>  $\oplus^{string\_agg}$  essentially is `||` (string concatenation) which is not commutative.

## Filtered and Unique Aggregates

---

```
SELECT agg(e) FILTER (p)  -- or: ... FILTER (WHERE p)
FROM ...
```

- **FILTER** clause alters aggregate semantics (see \*):

$$\begin{array}{l} \vdots \\ x \leftarrow e(t) \\ \text{if } x \neq \text{NULL} \wedge p(x): \\ \vdots \end{array}$$

```
SELECT agg(DISTINCT e)
FROM ...
```

- Aggregates distinct (non-NULL) values of expression *e*.  
(May use **ALL** to flag that duplicates are expected.)



## Evaluating Expressions on Minimal/Maximal Rows

---

In math, **arg max** finds the *argument*  $x_m$  that maximises function  $f$ :

$$x_m = \mathbf{arg\ max}_{x \in S} f(x) \quad \Leftrightarrow \quad f(x_m) = \mathbf{max}_{x \in S} f(x)$$

- In SQL, aggregate **arg\_max** (**arg\_min**) evaluates expression  $e_1$  for the row with maximal (minimal)  $e_2$ :

```
SELECT arg_max( $e_1$ ,  $e_2$ )  
FROM    ...
```

- **Q:** Convenient but not essential. Can you simulate **arg\_max** using only SQL constructs introduced so far?

## 11 : Forming Groups of Rows

---

Once **FROM** has generated row bindings, SQL clauses operate row-by-row. After **GROUP BY**: operate group-by-group:

<b>SELECT</b> $e_1, \dots, e_m$	-- 5
<b>FROM</b> ...	-- 1
<b>WHERE</b> ...	-- 2
<b>GROUP BY</b> $g_1, \dots, g_n$	-- 3
<b>HAVING</b> $p$	-- 4

- All rows that agree on all expressions  $g_i$  (the *set* of **grouping criteria**) form one **group**.
- $\Rightarrow$  Steps 4 and 5 process groups (*not* individual rows). This affects expressions  $p$  and the  $e_j$ .

## GROUP BY Partitions Rows

---

`SELECT ...`  
`FROM ...`  
`GROUP BY  $A_1$`   
`HAVING ...`

← evaluated once per **group** (not per row)

	$A_1$	$A_2$	...
	$\vdots$	$\vdots$	$\vdots$
the $x_i$ group {	$x_i$	$y_{i1}$	$\vdots$
	$x_i$	$y_{i2}$	$\vdots$
the $x_j$ group {	$x_j$	$y_{j1}$	$\vdots$
	$x_j$	$y_{j2}$	$\vdots$
	$\vdots$	$\vdots$	$\vdots$

Grouping **partitions** the row bindings:

- there are no empty groups
- each row belongs to exactly one group

## GROUP BY Changes Field Types From $\tau$ To $\text{bag}(\tau)$

---

✓	✗		**		✓
↓	↓		↓		↓
SELECT t.b, t.d	SELECT t.d		SELECT any_value(t.b) AS b, sum(t.d) AS "Σd"		
FROM T AS t	FROM T AS t		FROM T AS t		
GROUP BY t.b			GROUP BY t.b		

- $t.d$  references current group of  $d$  values: violates 1NF!  
 $\Rightarrow$  After **GROUP BY**: **must** use aggregates on field values.
- $t.b$  references current group of  $b$  values **all of which are equal** in a group  $\Rightarrow$  SQL: using just  $t.b$  is OK.
- ( $**$  aggregate  $\text{any\_value}(e)$  picks one arbitrary value among the equal  $e$  values.<sup>10</sup>)

<sup>10</sup> Indeed, in DuckDB:  $\text{any\_value}(e) \equiv \text{arbitrary}(e)$ .

## Aggregates Return One Result Per Group

---

```

SELECT t.b                                AS "group",
       count(*)                          AS size,
       sum(t.d)                          AS "Σd",
       bool_and(t.a % 2 = 0)             AS "∀even(a)",
       string_agg(t.a :: text, ';' )    AS "all a"
FROM   T AS t
GROUP BY t.b;

```

group	size	Σd	∀even(a)	all a
'x'	3	40	false	'1;3;5'
'y'	2	60	true	'2;4'

- **HAVING** *p* acts like **WHERE** but *after* grouping: *p* = false discards groups (not rows).

## Grouping Criteria

---

- The grouping criteria  $g_i$  form a set—order is irrelevant.
- Grouping on a **key** effectively puts each row in its own singleton group. (Typically a query smell 🐻.)
- Expressions  $e$  that are **functionally dependent** on the  $g_i$  are constant within a group (and thus *could* be used in `SELECT`).

Since SQL does not know about the FD, either

1. explicitly add  $e$  to the set of  $g_i$  (this will *not* affect the grouping), or
2. use `any_value(e)`.

## 12 : Bag and Set Operations

---

Tables contain **bags of rows**. SQL provides the common family of binary **bag operations** (*no row order*):

$q_1$	<b>UNION ALL</b>	$q_2$	-- $\cup^+$ (bag union)
$q_1$	<b>INTERSECT ALL</b>	$q_2$	-- $\cap^+$ (bag intersection)
$q_1$	<b>EXCEPT ALL</b>	$q_2$	-- $\setminus^+$ (bag difference)

- Row types (width, field types) of the  $q_i$  must match.
- With **ALL**, row multiplicities are respected: if row  $r$  occurs  $n_i$  times in  $q_i$ ,  $r$  will occur  $\max(n_1 - n_2, 0)$  times in  $q_1$  **EXCEPT ALL**  $q_2$  (**INTERSECT ALL**:  $\min(n_1, n_2)$ ).
  - Without **ALL**: obtain **set semantics** (no duplicates).

## 13 : Multi-Dimensional Data

---

- Relational representation of *measures* (*facts*) depending on multiple parameters (*dimensions*).
- Example: table `prehistoric` with *dimensions* `class`, `herbivore?`, `legs`, *fact* `species`:

<u>class</u>	<u>herbivore?</u>	<u>legs</u>	<u>species</u>
'mammalia'	true	2	'Megatherium'
'mammalia'	true	4	'Paraceratherium'
'mammalia'	false	2	NULL
'mammalia'	false	4	'Sabretooth'
'reptilia'	true	2	'Iguanodon'
'reptilia'	true	4	'Brachiosaurus'
'reptilia'	false	2	'Velociraptor'
'reptilia'	false	4	NULL

Table `prehistoric`



## Multiple GROUP BYs: GROUPING SETS

---

- Analyze (here: group, then aggregate) table  $T$  along multiple dimensions  $\Rightarrow$  perform separate GROUP BYs on each relevant dimension:
- SQL syntactic sugar:


```
SELECT  $e_1, \dots, e_m$ 
FROM    $T$ 
GROUP BY GROUPING SETS ( $G_1, \dots, G_n$ )
```

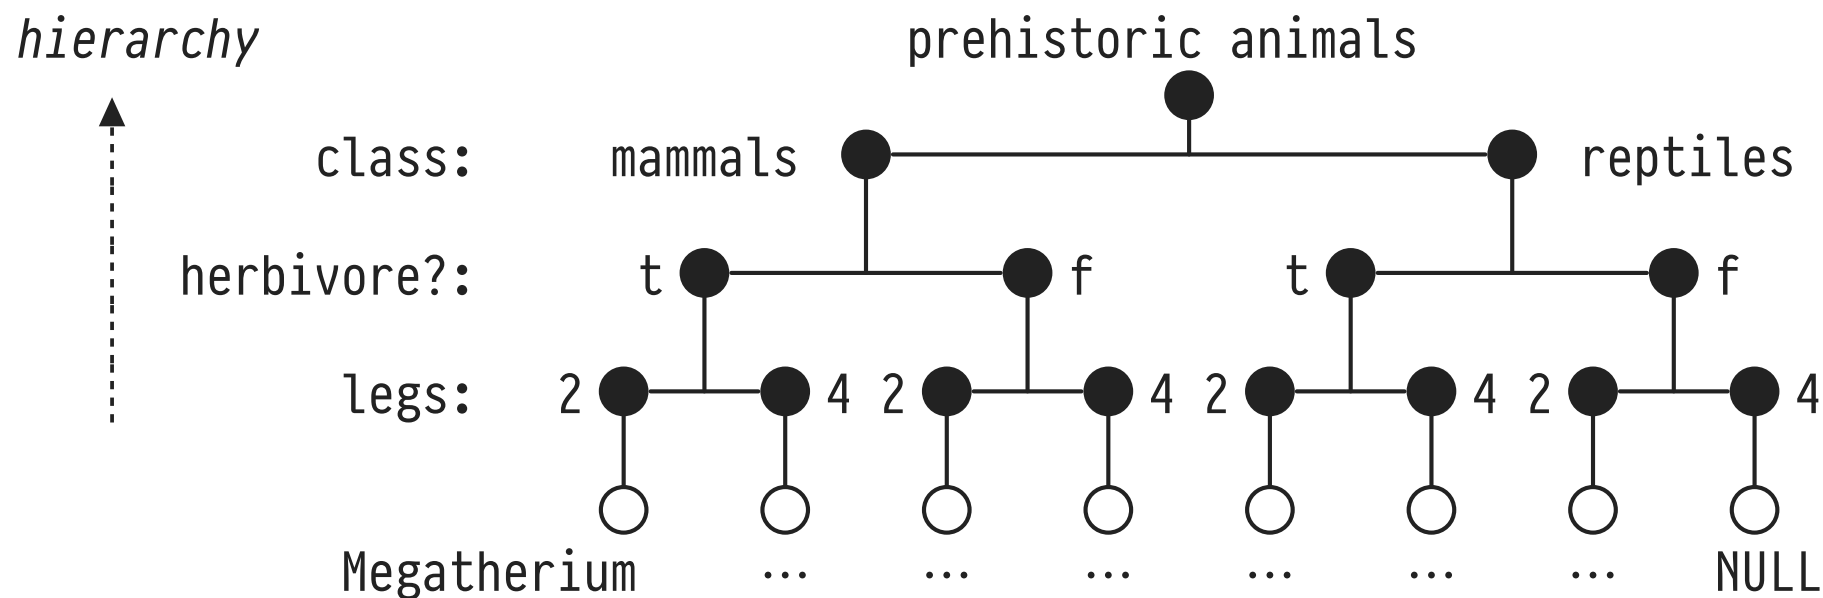
--  $G_i$ : grouping criteria  
-- sets in (...)

- Yields  $n$  GROUP BY queries  $q_i$ , glued together by UNION ALL. If non-aggregate  $e_j \notin G_i$ ,  $e_j \equiv \text{NULL}$  in  $q_i$ .

## Hierarchical Dimensions: **ROLLUP**

- **Group along a path** from any node  $G_n$  up to the root:


ROLLUP  $(G_1, \dots, G_n) \equiv$  GROUPING SETS  $((G_1, \dots, G_{n-1}, G_n)$   
 $, (G_1, \dots, G_{n-1}), \dots$   
 $, (G_1),$   
 $, ())$    $\downarrow$  *hierarchy*



## Analyze All Dimension Combinations: CUBE

---

●  $\equiv$  slice for herbivore? = true



$\text{CUBE } (G_1, \dots, G_n) \equiv \text{GROUPING SETS } ((G_1, \dots, G_n)$	$\left. \begin{array}{c} , \vdots \\ , ( ) \end{array} \right\}$	all $2^n$ subsets considered
--	--	------------------------------------

## 14 : SQL Evaluation vs. Reading Order

---

```

SELECT DISTINCT ON (es 7) es 3, aggs 6
FROM    qs      1
WHERE   p       2
GROUP BY es     4
HAVING  p       5

```

```


    UNION / EXCEPT / INTERSECT 8 } repeated 0 or more times,
    :                               } all evaluated before 9

```

```

ORDER BY es 9
OFFSET  n 10
LIMIT  n 10

```

- Reading order is: (7, 3, 6, 1 , 2, 4, 5, 8)<sup>+</sup>, 9, 10.
- (DuckDB's “friendly SQL”: admit FROM before SELECT.)

## Query Nesting and (Non-)Readability

---

```
SELECT ...  
FROM    (SELECT ...  
        FROM    (SELECT ...  
                FROM    ...  
                :    ) AS descriptive  
        :    ) AS ...  
:
```

- The more complex the query and the more useful the *descriptive* name becomes, the deeper it is buried. 🙄
- Query is a **syntactic monolith**. Tough to develop a query in stages/phases and assess the correctness of its parts.

## 15 | The `let...in` of SQL: `WITH` (Common Table Expressions)

---

Use **common table expressions (CTEs)** to bind table names *before* they are used, potentially multiple times:

```
WITH
  T1(c11, ..., c1, k1) AS [NOT MATERIALIZED] --
    (q1), --
    ⋮ --
  Tn(cn1, ..., cn, kn) AS [NOT MATERIALIZED] --
    (qn) --
q                                     -- q may refer to all tables Ti
```

} Query  $q_i$  may refer to tables  $T_1, \dots, T_{i-1}$

- “Literate SQL”: Reading and writing order coincide.
- Think of `let T1 = q1, ... Tn = qn in q` in your favorite FP language. The  $T_i$  are undefined outside `WITH`.

## SQL With **WITH** — Sample Uses

---

1. **Define queries in stages**, intermediate results in tables  $T_i$ . May use  $q \equiv \text{TABLE } T_i$  to debug stage  $i$ .
2. **MATERIALIZED**: **avoid recomputation** of common subqueries.<sup>11</sup>
3. **Bundle a query with test data**:


```
WITH
prehistoric(class,"herbivore?",legs,species) AS (
  VALUES ('mammalia',true,2,'Megatherium'),
         :
         ('reptilia',false,4,NULL)
)
SELECT max(p.legs)
FROM   prehistoric AS p
```

<sup>11</sup> DuckDB performs materialization by default. Modifier **NOT MATERIALIZED** will *inline*  $q_i$  at its use sites (if  $q_i$  is “simple” and inlining is semantically sound).

## 16 | 🔧 Use Case: **WITH** (Dinosaur Body Shapes)

---

Paleontology: **dinosaur body shape** (height/length ratio) and **form of locomotion** (using 2 or 4 legs) correlate:



- Use this correlation to infer bipedality (quadropedality) in incomplete dinosaur data sets:

<u>species</u>	height	length	legs
Gallimimus	2.4	5.5	?



## Dinosaur Body Shapes

---

<u>species</u>	height	length	legs
Ceratosaurus	4.0	6.1	2
Deinonychus	1.5	2.7	2
Microvenator	0.8	1.2	2
Plateosaurus	2.1	7.9	2
Spinosaurus	2.4	12.2	2
Tyrannosaurus	7.0	15.2	2
Velociraptor	0.6	1.8	2
Apatosaurus	2.2	22.9	4
Brachiosaurus	7.6	30.5	4
Diplodocus	3.6	27.1	4
Supersaurus	10.0	30.5	4
Albertosaurus	4.6	9.1	NULL
Argentinosaurus	10.7	36.6	NULL
Compsognathus	0.6	0.9	NULL
Gallimimus	2.4	5.5	NULL
Mamenchisaurus	5.3	21.0	NULL
Oviraptor	0.9	1.5	NULL
Ultrasaurus	8.1	30.5	NULL

Table dinosaurs

## Dinosaur Body Shapes

---

```
WITH  
bodies(legs, shape) AS (  
  SELECT d.legs, avg(d.height / d.length) AS shape  
  FROM   dinosaurs AS d  
  WHERE  d.legs IS NOT NULL  
  GROUP BY d.legs  
)  
:
```

<u>legs</u>	shape
2	0.447
4	0.201

Transient Table **bodies**

## Dinosaur Body Shapes

---

- **Query Plan:**<sup>12</sup>

0. Assume average body shapes in `bodies` are available
1. Iterate over all dinosaurs `d`:
  - If locomotion for `d` is known, output `d` as is
  - If locomotion for `d` is unknown:
    - Compute body shape for `d`
    - Find the `shape` entry `b` in `bodies` that matches `d`'s shape the closest
    - Use the locomotion (column `legs`) in `b` to complete `d`, output completed `d`

<sup>12</sup> In this course, *query plan* refers to a “plan of attack” for a query problem, not `EXPLAIN` output.