

# Advanced SQL

---

③

Standard and Non-Standard Data Types

Winter 2025/26

Torsten Grust  
Universität Tübingen, Germany

## 1 | Data Types in (DuckDB's) SQL

---

- The set of supported **data types** in DuckDB is varied:<sup>1</sup>

```
SELECT string_agg(t.type_name) AS "data types"  
FROM duckdb_types AS t  
WHERE t.database_name = 'memory';
```

data types  
varchar

bigint, binary, bit, blob, boolean, char, date, datetime,  
decimal, double, enum, float, hugeint, int, interval, list,  
map, numeric, real, row, struct, text, time, timestamp, uuid,  
⋮

---

<sup>1</sup> See [https://duckdb.org/docs/sql/data\\_types/overview](https://duckdb.org/docs/sql/data_types/overview)

## 2 | SQL Type Casts

Convert type of value  $e$  to  $\tau$  at runtime via a type cast:

<b>CAST</b> ( $e$ AS $\tau$ )	-- SQL standard
$e :: \tau$	-- shorthand, cf. FP
<b>TRY_CAST</b> ( $e$ AS $\tau$ )	-- yield NULL on failure

-  Type casts can fail at query runtime.
  - SQL performs **implicit casts** when the required target type is unambiguous (e.g. on insertion into a table column):

```
INSERT INTO T(a,b,c,d) VALUES (6.2, NULL, 'true', '0')
                                ↑      ↑      ↑      ↑
-- implicitly casts to:  int  text  boolean  int
```

## Literals (Casts From Type `text`)

---

SQL supports **literal syntax** for dozens of data types in terms of **casts from type `text`**:

<code>CAST('&lt;literal&gt;' AS <math>\tau</math>)</code>	}	succeeds if <code>&lt;literal&gt;</code> has a valid interpretation as $\tau$ (without cast $\Rightarrow$ type <u>text</u> )
<code>'&lt;literal&gt;' :: <math>\tau</math></code>		
$\tau$ <code>'&lt;literal&gt;'</code>		

- Embed complex literals (e.g., dates/times, JSON, enumerations) in SQL source.
- Casts from `text` to  $\tau$  attempted **implicitly** if target type  $\tau$  known. Also vital when importing data from text or CSV files (*input conversion*).

### 3 | Text Data Types

<u>text</u>	-- UTF-8 string of unlimited length
<u>string</u>	--
<u>char</u>	--
<u>varchar</u>	-- } synonyms for type text

- Text values are of unlimited length. Specifying a “maximum length” (as in `char(n)`, `varchar(n)`) is accepted for SQL compatibility but has no effect on the system or storage.
  -  One UTF-8 character may occupy more than one byte.
- **NB.** SQL text literals are enclosed in *single* quotes '`...`' (or `$<id>$...$<id>$` with matching `<id>`).

## 4 | NUMERIC:<sup>2</sup> Large Fixed-Point Decimals

numeric(*width*, *scale*)

scale

1234567.890

width (# of digits)


- Shorthands: `numeric(width, 0) ≡ numeric(width)`, `numeric ≡ numeric(18, 3)`.
- Exact arithmetics, but may be computationally heavy for large *width* values. Internal representation are integers:

<i>width</i>	<b>internal</b>	<b>size (bytes)</b>
1-4	int16	2
5-9	int32	4
10-18	int64	8
19-38	int128	16

<sup>2</sup> Synonymous: `decimal`.

## 5 | Timestamps and Time Intervals

---



- Types: `timestamp`  $\equiv$  `(date, time)`. Casts between types: `timestamp`  $\rightarrow$  `time/date` ✓, `date`  $\rightarrow$  `timestamp` assumes 00:00:00. With time zone support: `{time,timestamp}` with time zone.<sup>3</sup>
- Type `interval` represents timestamp differences.
- Resolution: `timestamp/time/interval`: 1  $\mu$ s, `date`: 1 day.

<sup>3</sup> Type name shorthands: `timetz`, `timestamptz`.

## Date/Time/Interval Literals: DuckDB (ISO 8601 Format)

---

- Date literals (format YYYY-MM-DD): `date '1968-08-26'`,  
`'1904-05-30' :: date.`
- Time literals (hh:mm:ss[.zzzzzz][+-TT[:tt]]):  
`'14:20:33.982+00' :: time.`
- timestamp literal  $\equiv$  '`<date literal>_<time literal>`'
- interval literal (fields optional, `s` may be fractional)  $\equiv$   
`'<n>years <n>months <n>days <n>hours <n>mins <s>secs'`
- Built-in date/time/interval literals:
  - `date: today()`
  - `timestamp: now(), today() :: timestamp`
  - `interval: interval(<n>) <unit>` (… `interval(6) hours` …)

## Computing with Time

---

- Timestamp arithmetic via `+`, `-` (also `*`, `/` with intervals):

```
SELECT now() :: timestamp - today() :: timestamp;
```

interval
12:58:58.022

- DuckDB: *Extensive* library of date/time/timestamp functions including:
  - `strftime()`, `strptime()` (flexible formatting/parsing)
  - `age()`, `month_name()`, `extract(<part> from •)`
  - `make_date(•,•,•)`, `make_time(•)`, `make_timestamp(•)`
  - comparisons (`=`, `<`, ...)

## 6 | Enumerations

---

Create a *new* type  $\tau$ , incomparable with any other. Explicitly **enumerate** the literals  $v_i$  of  $\tau$ :

---

```
CREATE TYPE  $\tau$  AS ENUM ( $v_1, \dots, v_n$ );
```

---

```
SELECT  $v_i::\tau$ ;
```

- Literals  $v_i$  in case-sensitive quoted notation ' $\dots$ '.  
(Compact storage in  $\lceil \log_2(n)/8 \rceil$  bytes, regardless of literal length.)
- Implicit ordering:  $v_i < v_{i+1}$  (aggregates `MIN`, `MAX` ✓).

## 7 | Bit Strings

---

- Data type `bit` stores strings of binary digits (storage: 1 byte per 8 bits + constant small overhead).
- Literals:

```
SELECT '0'::bit, '00101010'::bit, bitstring('1010',12)
--                                         ^
--                                         ≡ 000000001010 (zero-pad to 12 bits)
```

- Bitwise operations: `&` (and), `|` (or), `xor`, `~` (not), `<>/>>` (shift left/right), `get_bit(·,·)`, `set_bit(·,·,·)`, ...
- Operation on bit strings: `bit_count(·)` (counts 1-bits), `bit_length(·)`, `octet_length(·)` (number of bytes), ...
- Aggregates: `bit_and(·)`, `bit_or(·)`, `bitstring_agg(·,·,·)`, ...

## 8 | Binary Arrays (BLOBs)

---

Store **binary large object blocks** (BLOBs; ,  in column  $B$  of type `blob`) inline with alpha-numeric data. BLOBs remain *uninterpreted* by DBMS:


$\dots$	$K$	$B :: \text{blob}$	$P$	$\dots$
	$\vdots$		$\vdots$	
	$k_i$		$p_i$	
	$k_j$		$p_j$	
	$\vdots$	$\vdots$	$\vdots$	

Table T

- Typical setup:
  - BLOBs stored alongside identifying **key** data (column  $K$ ).
  - Additional **properties** (meta data, column(s)  $P$ ) made explicit to filter/group/order BLOBs.

## Encoding/Decoding BLOBs

- Import **blob** data via textual encoding (e.g., base64) or directly from binary files via `read_blob(•)`:



⚠ Maximum size of a **blob** column value is 4 GB.

## 9 | JSON (JavaScript Object Notation)

---

**JSON** defines a textual data interchange format. Easy for humans to write and machines to parse (see [json.org](http://json.org) ):

---

```
<object> ::= {} | { members }
<members> ::= pair | pair , members
<pair> ::= string : value
<array> ::= [] | [ elements ]
<elements> ::= value | value , elements
<value> ::= string | number | true | false | null
          | array | object
```

---

- SQL:2016 defines SQL $\leftrightarrow$ JSON interoperability. JSON *<value>*s may be constructed/traversed and held in table cells (we still consider 1NF to be intact).

## JSON Sample <value>s

```
        <members>
        { "title":"The Last Jedi", "episode":8 }
        ↑
<object>           <pair>
```

## Table T (see Chapter 02):

```
<elements> { [ { "a":1, "b":"x", "c":true, "d":10 },  
  { "a":2, "b":"y", "c":true, "d":40 },  
  { "a":3, "b":"x", "c":false, "d":30 },  
  { "a":4, "b":"y", "c":false, "d":20 },  
  { "a":5, "b":"x", "c":true, "d":null } ]  
  ↑  
  <number>           <array> (of <object>s)
```

## JSON in DuckDB: Type `json`

---

Quoted literal syntax embeds JSON values in SQL queries.  
Wrap in function `json(•)` to validate and minify JSON syntax:

```
VALUES (1, json('{"b":1, "a":2}')),  
        (2, json('{"a":1, "b":2, "a":3}')),  
        (3, json('[ 0,    false,null ]'));
```

col0	col1 (:: json)
1	{"b":1,"a":2}
2	{"a":1,"b":2,"a":3}
3	[0,false,null]

## Navigating JSON Values

---

- Given a JSON value `value`, access object field `f` (or array element at index `i`) using binary operator `->:4`

---

`value->'f'` } yields a JSON value, navigate further or  
`value->i` } cast to atomic type for further computation


---

- Path navigation:
  - Chain navigation steps: `value->'f1'->'f2'->...->'fn'`.
  - Use JSONPath expression: `value->'<path>'`, see next slide.

<sup>4</sup> Accessing non-existent `f/i` yields `NULL`. **NB:** Array indexes `i` are 0-based.

## Navigating JSON Values: JSONPath Expressions

- JSON values describe **tree-shaped** data:

$$o \equiv \{ "a": 0, "b": \{ "c": [1,2], "d": 3 \} \}$$



- Navigate from root \$, return (array of) json values:

```
SELECT o->'<path>';
```

--

-- embed JSONPath expression as quoted text literal

# Navigating JSON Values: The JSONPath Language



**SELECT o->'<path>' AS j**

**SELECT json\_extract(o, '<path>') AS j**

<path>	result j	
\$	v <sub>0</sub>	root
\$.*	[v <sub>1</sub> , v <sub>2</sub> ]	all child values of the root <sup>5</sup>
\$.a	v <sub>1</sub>	child a of the root
\$.b.d	v <sub>4</sub>	grandchild d below child b
\$.b.c[1]	v <sub>6</sub>	2nd array element of array c
\$.b.c[*]	[v <sub>5</sub> , v <sub>6</sub> ]	all array elements in array c
\$.b.c[#-2]	v <sub>5</sub>	second to last element in array c
\$.b.**	[v <sub>3</sub> , v <sub>4</sub> , v <sub>5</sub> , v <sub>6</sub> ]	all nodes in subtree below b

<sup>5</sup> In this case, `->` returns a *SQL array* of values (type `json[]`). Chapter 04 will cover SQL arrays.

## Bridging between JSON Objects and SQL Rows

---

- A JSON object `o` may be casted into a row value `t` provided that `o`'s **structure is known**:

```
-- o
SELECT json('{"a":1,"b":true}') :: row(a int, b boolean);
-- t
```

- Likewise, `json(•)` turns row values `t` into JSON objects:

```
SELECT json(t) AS o
FROM T AS t;
```

```
o (:: json)
{"a":1,"b":"x","c":true,"d":10}
:
{"a":5,"b":"x","c":true,"d":null}
```

## Constructing JSON Values From a Table of Inputs

- `json_group_array(v)`: collect values  $v$  in a JSON array:

---

```
SELECT json_group_array(t.v) AS o
FROM  (VALUES (v1),
              (v2),
              (v3)) AS t(v);
```

---

[  $v_1$ ,  
 $v_2$ ,  
 $v_3$  ]

- `json_group_object(k,v)`: construct JSON object  $o$  from a table of  $(k,v)$  pairs:

---

```
SELECT json_group_object(t.k, t.v) AS o
FROM  (VALUES (k1, v1),
              (k2, v2),
              (k3, v3)) AS t(k,v);
```

---

{  $k_1: v_1$ ,  
 $k_2: v_2$ ,  
 $k_3: v_3$  }

## 10 | Sequences

---

**Sequences** represent counters of type **bigint** ( $-2^{63} \dots 2^{63}-1$ ). Typically used to implement row identity/name generators:

---

<b>CREATE SEQUENCE</b> <i>seq</i>	-- sequence name
[ <b>INCREMENT</b> <i>inc</i> ]	-- advance by <i>inc</i> (default: $1 \equiv \uparrow$ )
[ <b>MINVALUE</b> <i>min</i> ]	-- range of valid counter values
[ <b>MAXVALUE</b> <i>max</i> ]	-- (defaults: $[1 \dots 2^{63}-1]$ )
[ <b>START</b> <i>start</i> ]	-- start (default: <i>min</i> if $\uparrow$ , <i>max</i> if $\downarrow$ )
[ <b>[NO] CYCLE</b> ]	-- wrap around or error( $\equiv$ default)?

---

- Columns can be tied to a sequence for key generation:

---

<b>CREATE SEQUENCE</b> <i>T_keys</i> ;
<b>CREATE TABLE</b> <i>T</i> ( <i>k</i> <b>int</b> <b>DEFAULT</b> <b>nextval</b> ('T_keys'), ...);

---

## Advancing and Inspecting Sequence State

---

- Counter state can be advanced (`nextval(•)`) and inspected (`currval(•)`):

```
-- seq: 41 → 42 → 1 → 2 → ... 41 → 42 → 1 → ...
CREATE SEQUENCE seq START 41 MAXVALUE 42 CYCLE;
:
SELECT nextval('seq');           --  ⇒ 41
SELECT nextval('seq');           --  ⇒ 42
SELECT currval('seq');           --  ⇒ 42
SELECT nextval('seq');           --  ⇒ 1   (wrap-around)
                                ↑
```



sequence/table names are not first class in SQL

- `c τ DEFAULT e` evaluates expression `e` (of type `τ`) whenever row insertion omits a value for column `c`.